A. Additional evaluation results

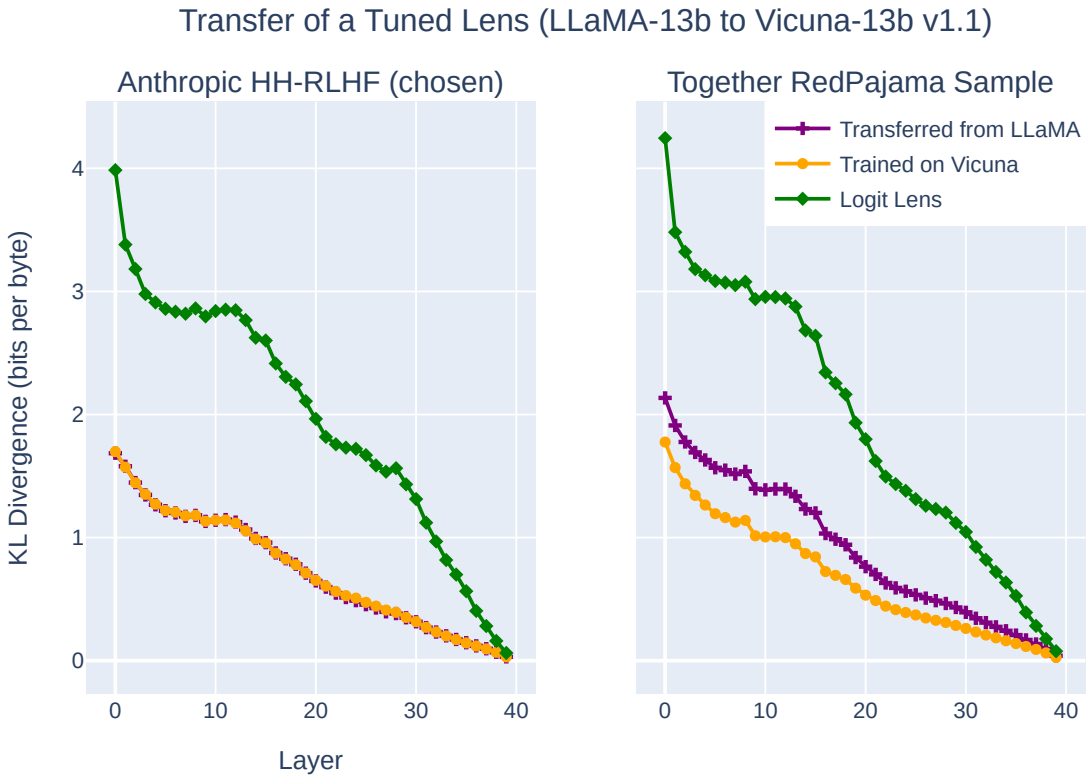

Figure 13. Comparing the performance of a lens specifically trained on Vicuna 13B vs a lens transferred from LLaMA 13B vs Vicuna 13B's Logit lens. Specifically, we messure the KL divergence between the lens at a specific layer and the model's final output.

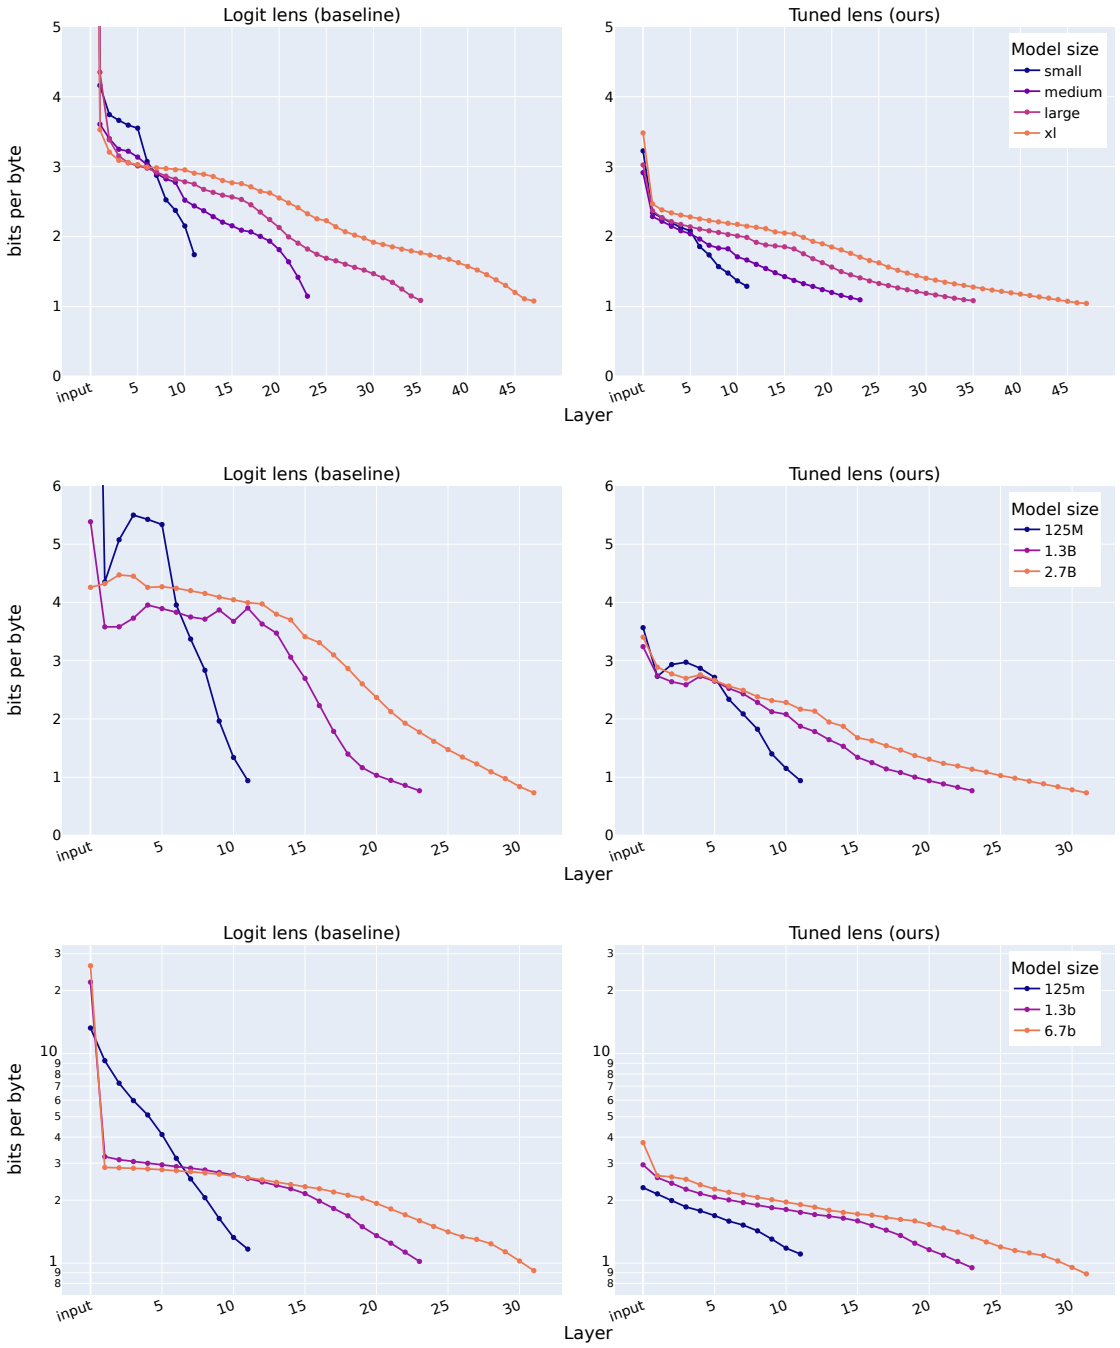

Figure 14. Perplexities of predictions elicited from OpenAI's GPT-2 models (top), EleutherAI's GPT-Neo models (middle), and Meta's OPT models (bottom). We omit OPT 350M as it uses a post-LN architecture.

B. Qualitative Results

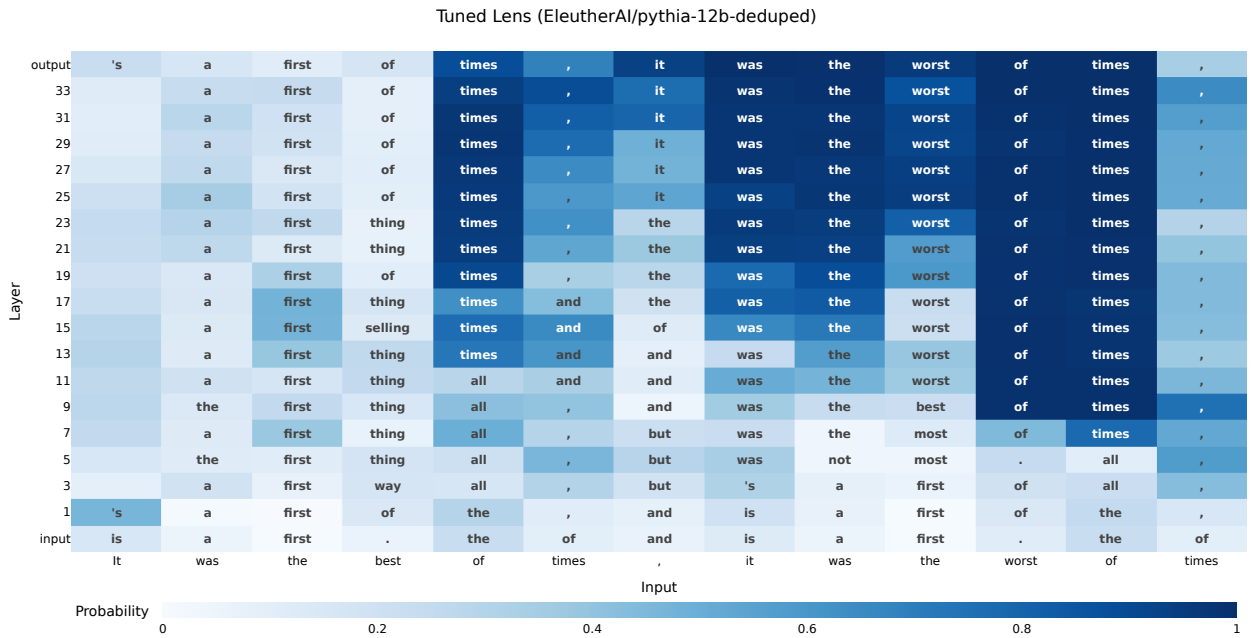

Figure 15. Tuned lens prediction trajectory for Pythia 12B on the first several words of Charles Dickens' "A Tale of Two Cities." Note that the latent prediction becomes very confident at early layers after the prefix "It was the best of times" has been processed, suggesting some degree of memorization.

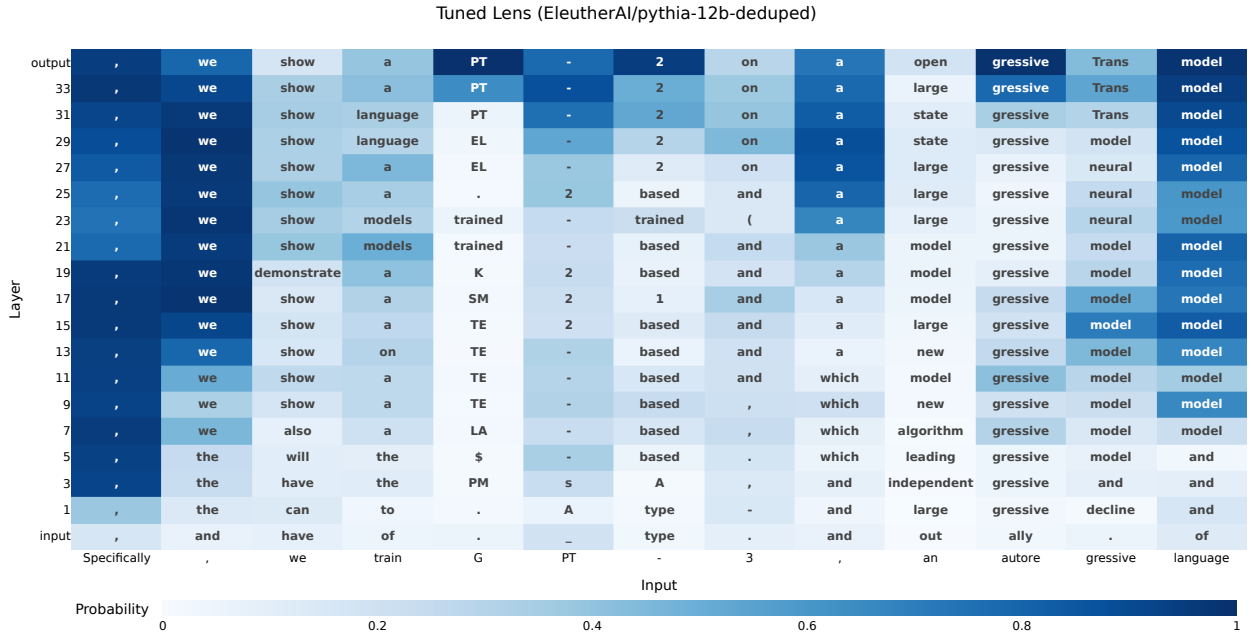

Figure 16. Tuned lens prediction trajectory for Pythia 12B prompted with the abstract of Brown et al. (2020).